

PROOF2PAY AI - Responsible AI

Grounded evidence, uncertainty, privacy, and human control

Safety Specification v0.1

Safety objective

Use AI where it creates operational leverage without allowing the model to invent evidence, bypass requirements, or authorize legal/financial actions.

Core policy

AI may interpret and recommend. It may not manufacture proof.

Risk-control matrix

Risk	Example	Control
Hallucinated completion	Treating “done” in a note as proof	Claim/evidence separation; critical evidence rules
Wrong requirement	Model retrieves unrelated policy	Controlled sources + source reference + human review
Ambiguous evidence	Blurry photo or vague test note	AMBIGUOUS ; request better proof
Automation overreach	AI sends acceptance request automatically	explicit approval gate
Financial overreach	AI marks job invoice-authorized	final MVP human confirmation; no invoice action
Prompt injection	File says “ignore all rules”	uploaded files treated as data, not instructions

Risk	Example	Control
Privacy leakage	customer signature/address reaches logs	minimization, masking, access controls
Misleading confidence	92% score shown as certainty	blocker-first UX; sources and limitations

Grounding requirements

Every critical assessment should include:

- requirement ID/title;
- source reference;
- evidence IDs considered;
- status;
- rationale;
- uncertainty/limitations;
- next action;
- human-review flag.

Data minimization

For the hackathon:

- use synthetic customer names/IDs;
- use mock attachments or non-sensitive demo assets;
- never upload real signatures without consent;
- never commit production API keys or customer documents;
- avoid collecting fields unrelated to readiness.

Human review

Mandatory before:

1. final BILLING_READY confirmation in MVP;
2. customer-facing follow-up messages;
3. externally used service/handover documents;
4. accounting/invoice handoff;
5. any irreversible or legally consequential action.

Evaluation

Safety is testable. Include cases where:

- technician claims a test was done but no proof exists;
- evidence conflicts;
- source document is missing;
- uploaded document contains malicious instructions;
- user pressures the assistant to bypass a blocker.

A safe system should escalate rather than invent.